

CS 482: Computational Techniques in Biological Sequence Analysis Homework #2

ERIC HAORAN HUANG¹

¹e48huang@uwaterloo.ca, 20880126, e48huang

Abstract

This assignment was an exercise in phylogeny, alignment-free methods and genome assembly.

Setup

I used Python 3.10.12 on the student server machines in this run with external dependencies described in 'requirements.txt'. Run 'pip install requirements.txt' for proper setup.

Part 1: K-mer Composition

All code and examples are found under 'q1/'. This problem was tackled in a $O(nk)$ runtime where we had the following strategy:

1. Grab the sequence using the BioPython.
2. Calculate the k-mer frequency array by iterating over all k-substrings (running $n - k + 1$ times with the length of the sequence being n).
 - ◇ For each k-mer, we generate all possible k-mers, e.g.: 'NA' will produce 'AA', 'CA', 'GA', 'TA'.
 - ◇ We do this by iterating through each k-mer string one nucleotide at a time and keeping all possible k-mer prefixes. We extend each prefix by the possible next nucleotide base according to IUPAC notation. This process takes $O(nk)$ times because we have $O(n)$ possible k-mers which we spend $O(k)$ time reconstructing all possible k-mers.
 - ◇ Taking the total k-mers possible, weigh each possible string equally across each k-mer. Add this weight to a k-mer frequency array whose index is defined as mapping of the possible kmer string to its lexicographic index in the 4^k frequency matrix. This can be easily found by setting the weights of A to 0, C to 1, G to 2, T to 3 and then turning the string (alias of base 4) to base 10.
3. Return and print out to a file as needed.

Run the code with the following line: 'python kmer_comp.py -i <input_file> -k <kmer-length>' with optional flags of '-o <output_dir>' to output to the file under '<output_dir>/<input_file_name>_len_<kmer_length>_k-mers.txt' and '--debug' for more logging information.

I decided to use this method of averaging over the possible k-mers for ambiguous bases, i.e. if we have 'NA' we give 0.25 frequency weights to 'AA', 'CA', 'GA', 'TA'. This assumes that given an ambiguous base, that there is equal chance of any base that represents it to be the true base. This might not necessarily be true, but is the best way to utilize the heuristic of the ambiguous base given.

Part 2: deBruijn Graph Construction

All code and examples are found under 'q2/'. This problem was tackled in a $O(nk \log n)$ runtime with the following strategy, where n is the number of strings, k is the length of the strings:

1. Calculate the reverse complements of all the strings, taking $O(nk)$ time.
2. Calculate the edge list, taking $O(nk)$ time of the joint set of the given set and the reverse complements.

3. Return the sorted list, taking $O(k)$ time in comparison and needing $O(n \log n)$ comparisons, or done in $O(nk \log n)$ time.

Run the code with the following line: `python build_deBruijn.py -i <input_file>` with optional flags of `-o <output_dir>` to output to the file under `<output_dir>/<input_file_name>_deBruijn.txt` and `--debug` for more logging information.

Part 3: Alignment-free analysis of viral phylogenies

All code and examples are found under 'q3/'. This problem heavily utilized the existing libraries as follows:

1. Downloaded the viral sequences using BioPython's Entrez package. This resulted in '.gb' files which were then cached and parsed through using BioPython again.
2. Defined metrics as following:
 - ◊ Euclidean distance. This metric is already a distance, no changes needed, used the L2-norm.
 - ◊ Cosine similarity. This metric is not a distance and in fact, given that each component itself must be non-negative (given that frequency must be greater than or equal to 0), according to this Wikipedia article*, it is bounded by $[0, 1]$, where it being 0 means that it is completely dissimilar, and being 1 means that it is an exact match. So, we need to translate this to distance which is bounded by $[0, \infty)$. Therefore, we need to map the 1 to a 0 and the 0 to an infinity. A natural solution then is given d as our cosine similarity, to take $\frac{1}{d} - 1$ as this maps 0 to infinity, and also maps 1 to 0. I came up with this by noting that $\lim_{d \rightarrow 0} \frac{1}{d} = \infty$ but we need to shift it by 1 for $\frac{1}{d} - 1$ to be 0 at $d = 1$.
 - ◊ Pearson Correlation. Similarly, I noted that according to Wikipedia*, that the expression will always lie between -1 and 1 , with -1 being completely negatively correlated and 1 being completely positively correlated. This might look like if given two vectors, the AA frequency increasing from one sequence to another will cause the frequency of CC to decrease. So, a negative correlation means that they are more dissimilar, whereas a positive correlation means that they are more similar. Therefore, we need to map -1 to ∞ and 1 to 0. Then we have, given the Pearson correlation of ρ :

$$\lim_{\rho \rightarrow -1} \frac{1}{\rho + 1} = \infty, \frac{1}{\rho + 1}|_{\rho=1} = \frac{1}{2}$$

which matches the proper value for $\rho = -1$, however is wrong for $\rho = 1$. Then, we can do:

$$\begin{aligned} \frac{2}{\rho + 1} - 1 &= \frac{2 - \rho - 1}{\rho + 1} \\ &= \frac{1 - \rho}{1 + \rho} \end{aligned}$$

which is what we implemented.

3. Used each of these distances on the kmer frequency vector calculated using q1, which created a distance matrix.
4. Took this distance matrix and used the DistanceTreeConstructor module to build a UPGMA tree.

Run the code with the following line: `python build_newick_tree.py -i <input_file> -c <input_candidate_file>` with optional flags of `-o <output_dir>` to output to that directory and `--debug` for more logging information. There is also a `-k` flag to attempt this with different k-mer lengths.

Then, using this result we can answer the given questions.

Subtype Clustering

The subtypes are very well clustered in the initial tree as shown in the following figures representing the Newick tree, generated through IcyTree. We can see that for every single tree, we have that the subtypes are all grouped together before grouping with any other subgroup.

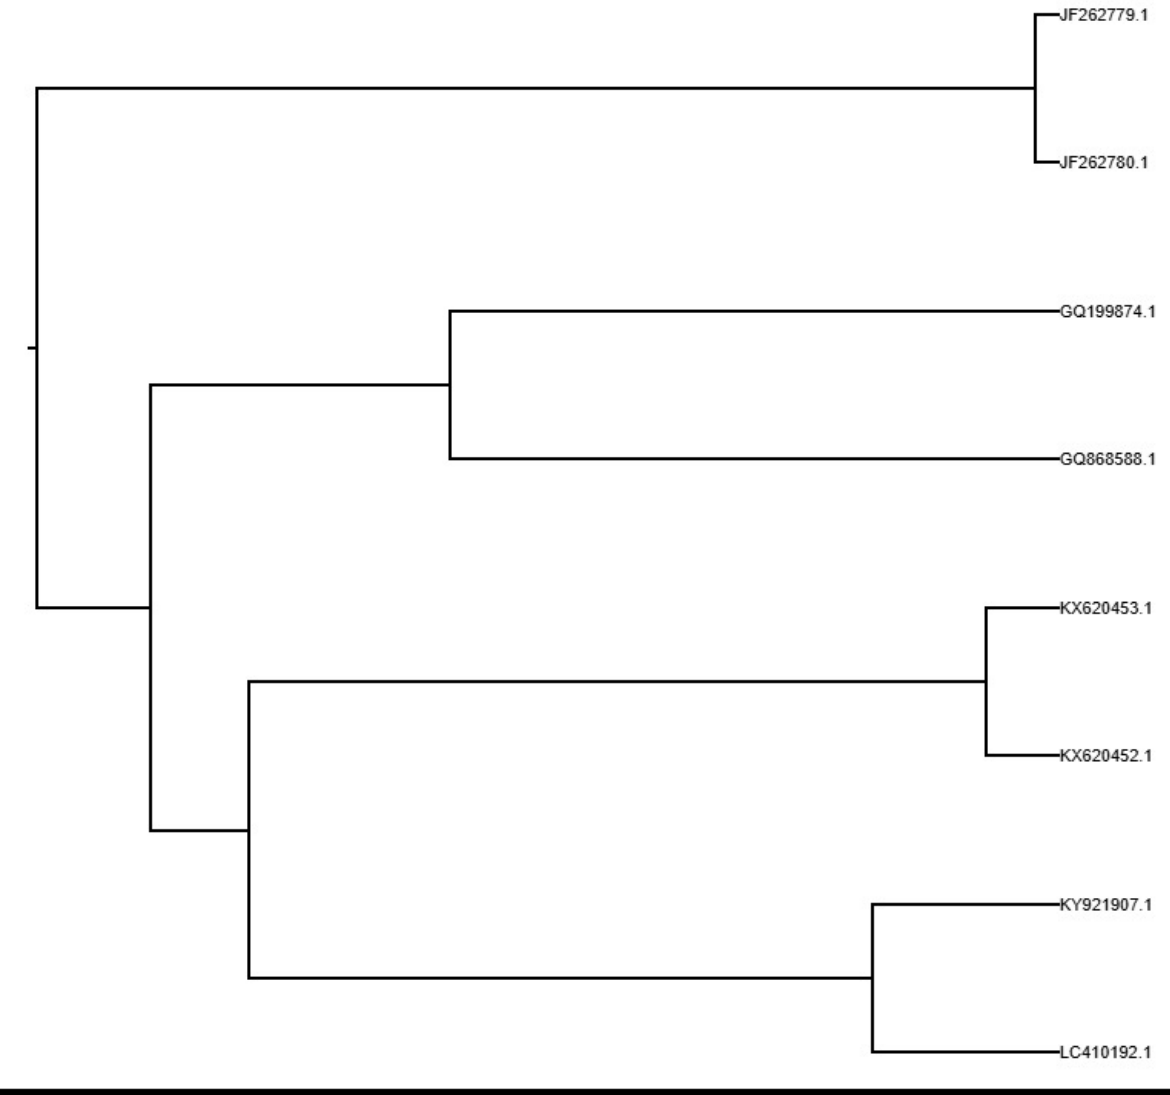


Figure 1: Euclidean Distance Newick Tree of the original subtypes

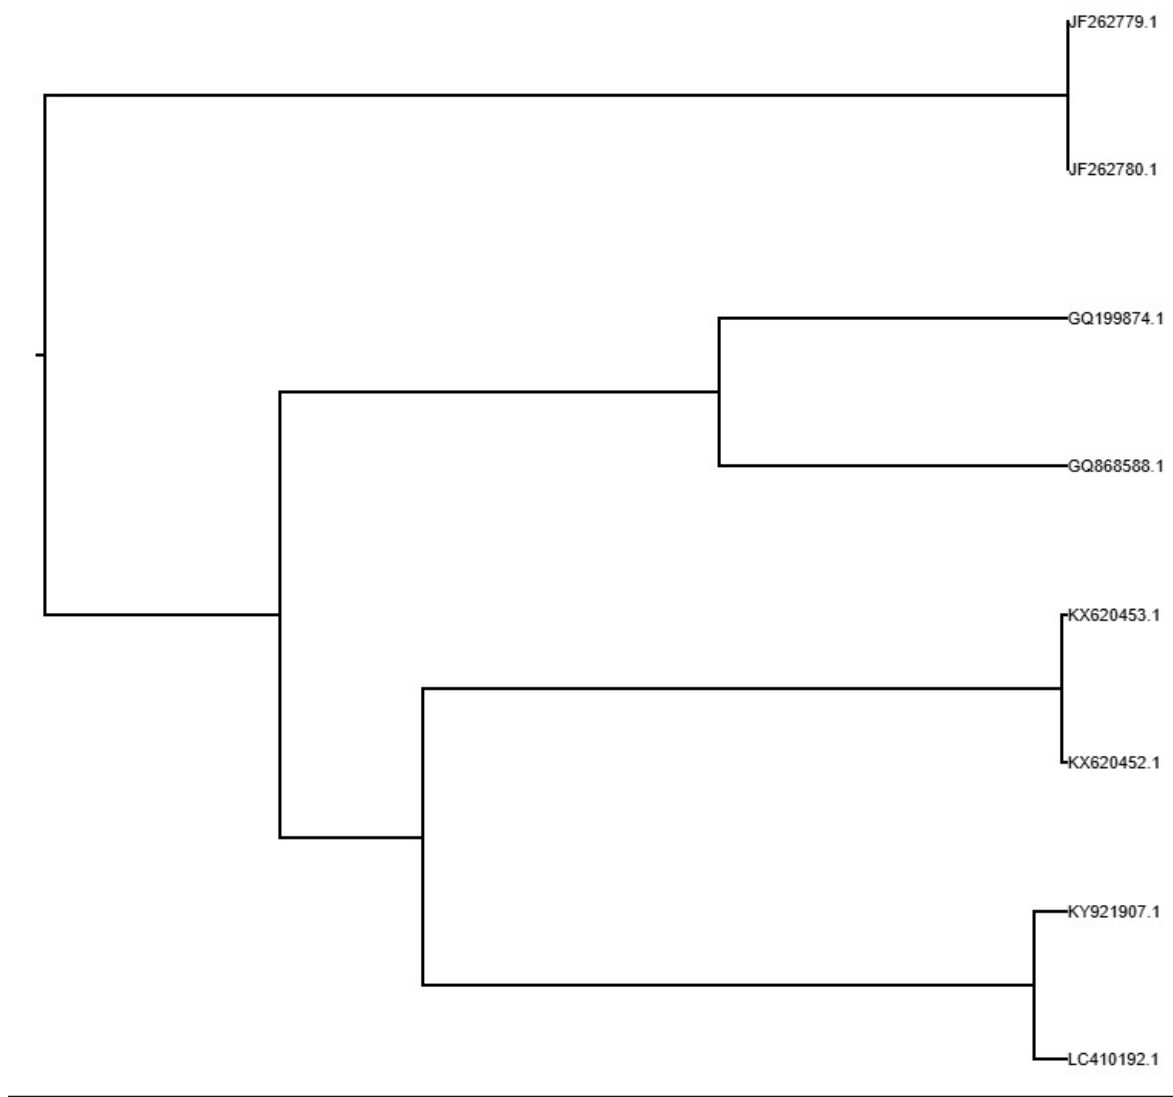


Figure 2: Cosine Similarity Newick Tree of the original subtypes

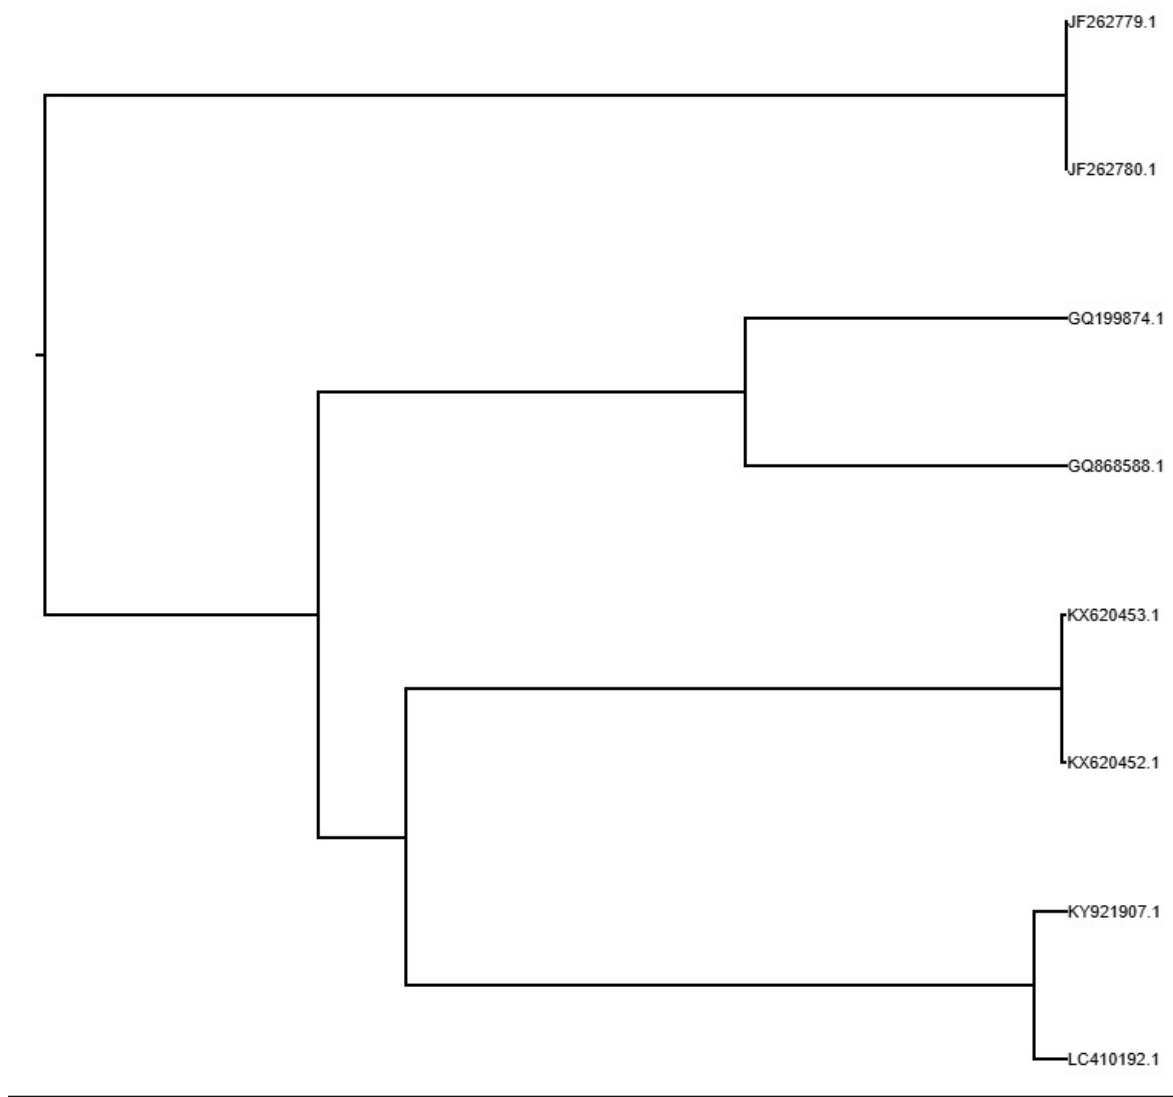


Figure 3: Pearson Correlation of Newick Tree of the original subtypes

Test Sequence

With the test sequence, we see a slight change in the tree, where we simply add sequence 9 in.

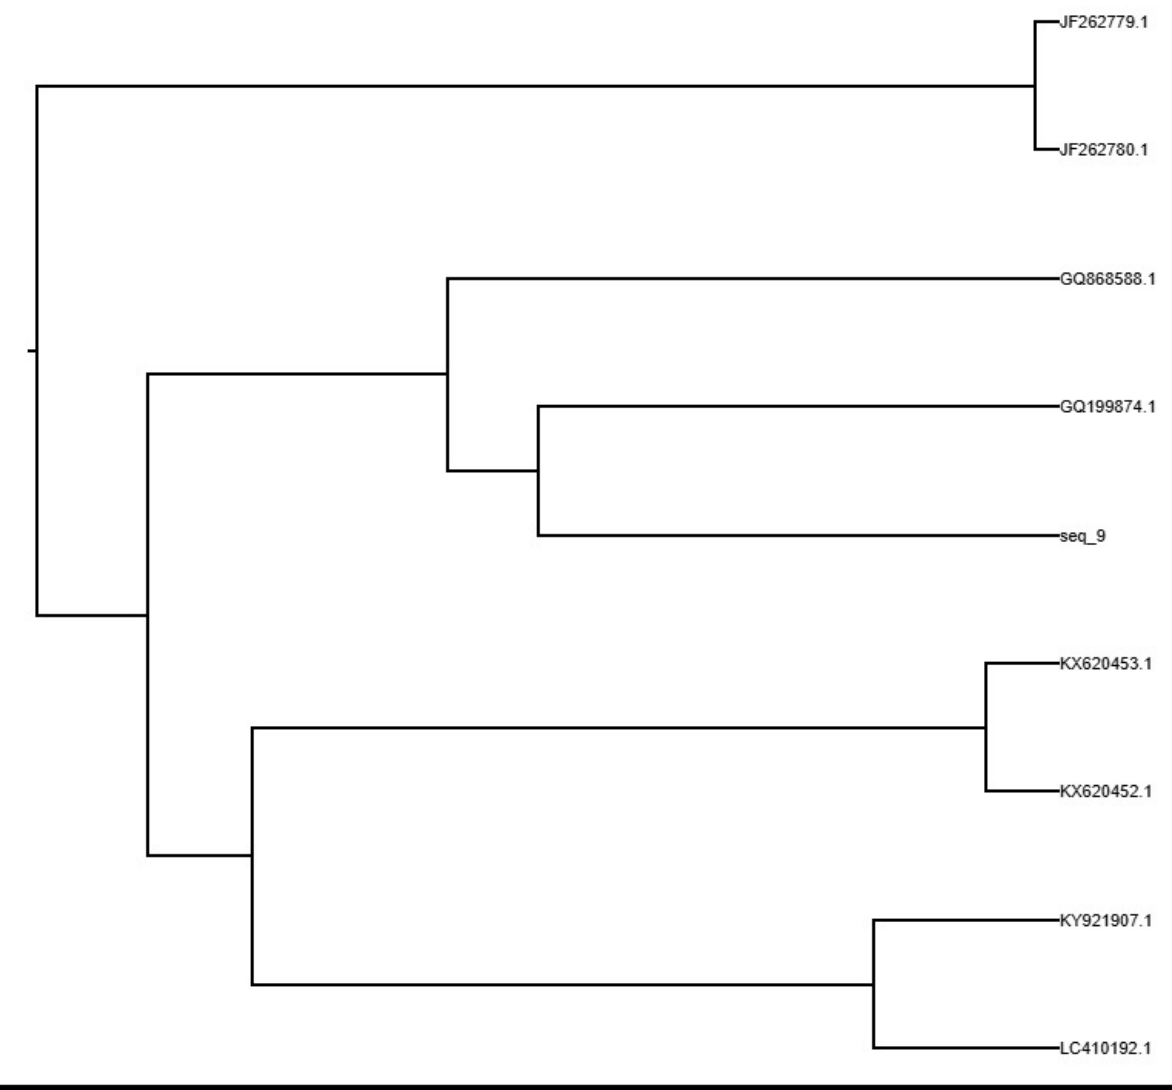


Figure 4: Euclidean Distance Newick Tree with the test sample

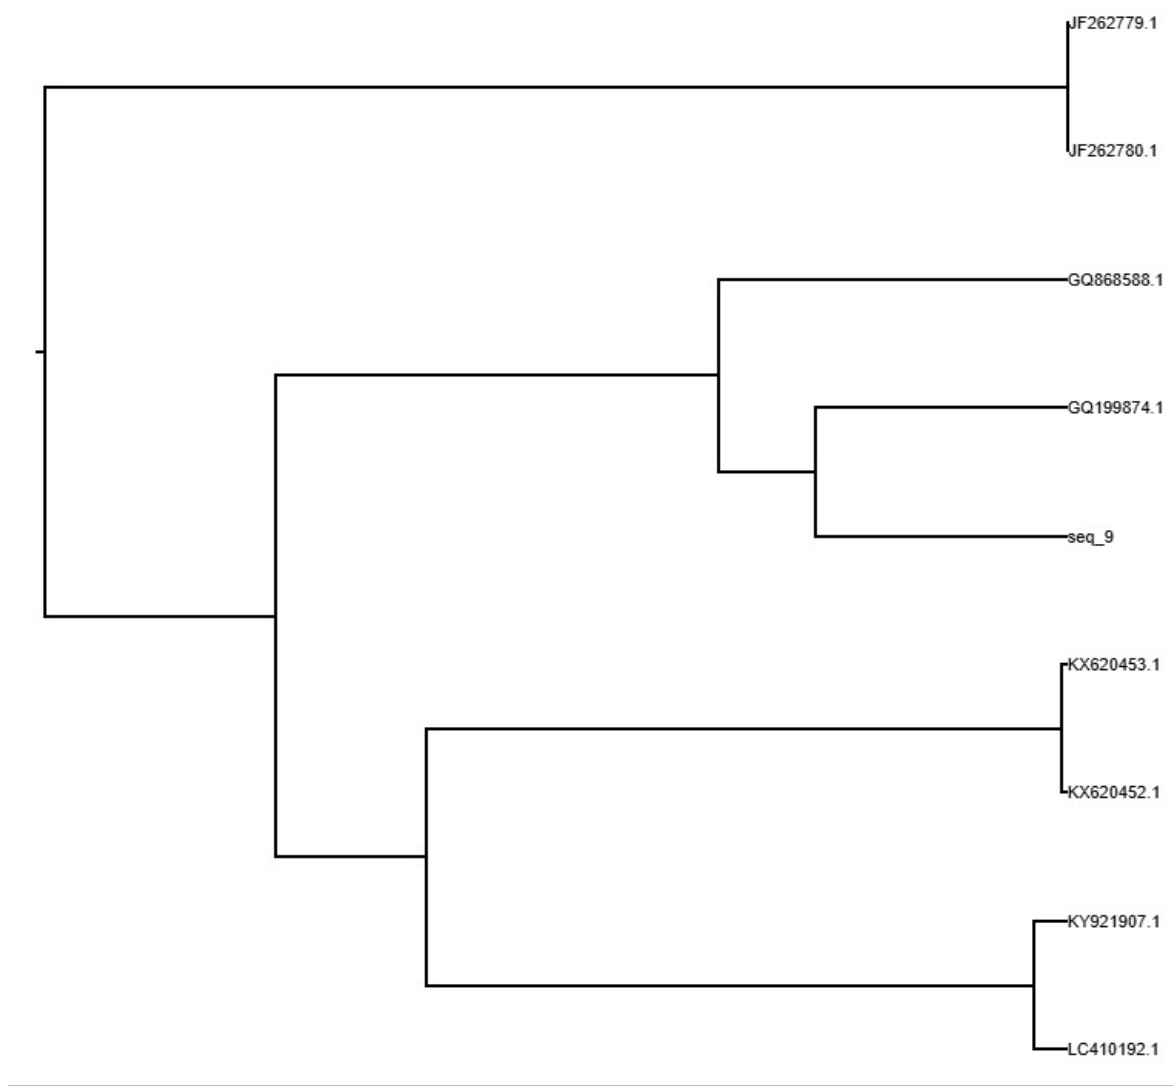


Figure 5: Cosine Similarity Newick Tree with the test sample

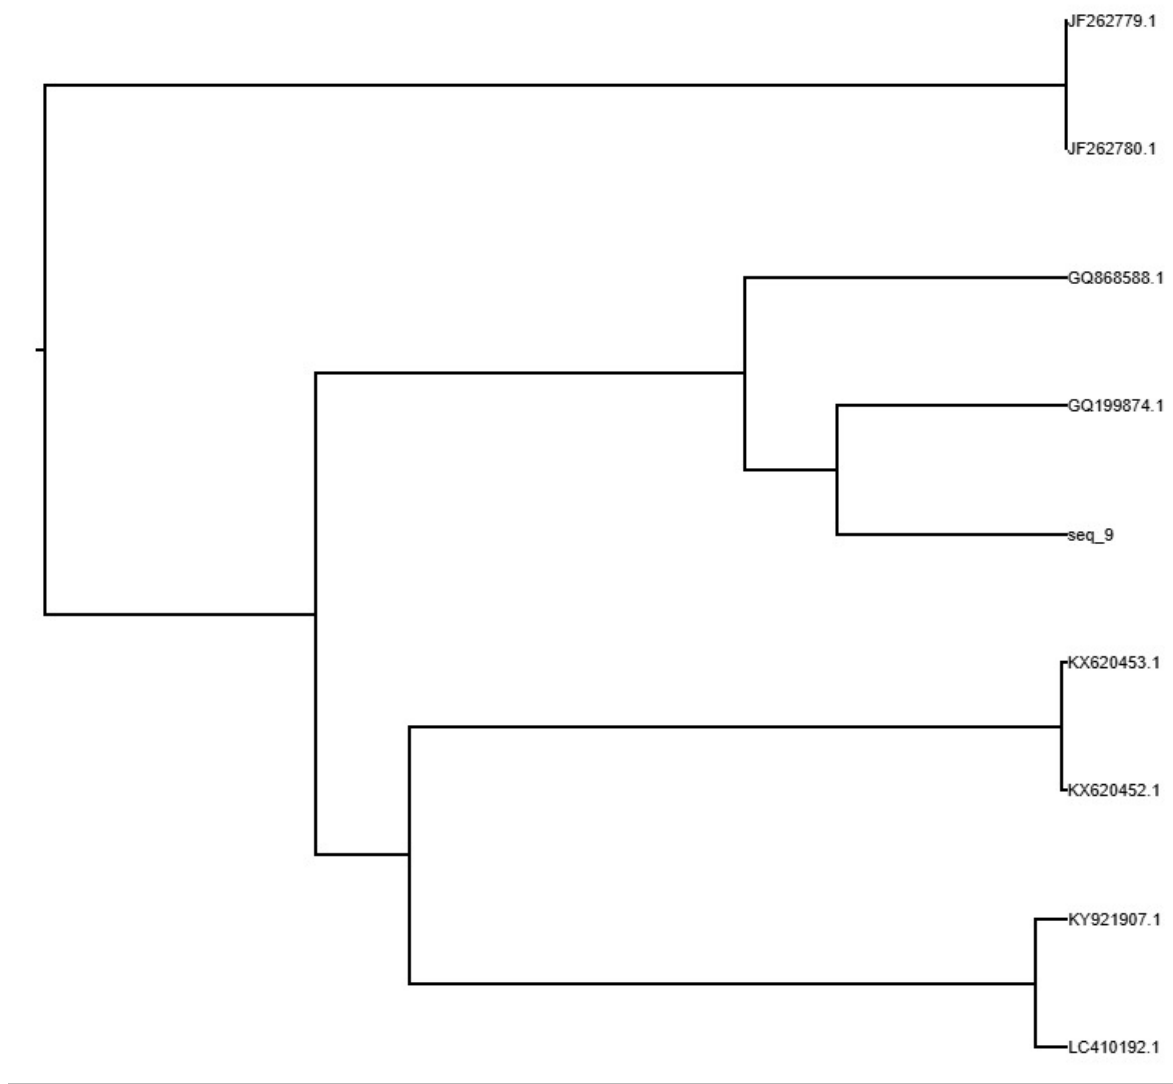


Figure 6: Pearson Correlation of Newick Tree with the test sample

Predicted subtype

In all cases, we have that our subtype sequence 9 belongs to subtype2, and is a sibling of GQ199874.1.

Appropriate Metric

In my opinion, I believe that cosine similarity to be the best metric as it is the most interpretable and grounded in what we want.

For Euclidean distance, a problem arises if we have two sequence frequencies of $[0, 1]$ and $[0, 2]$. If the direction is the same for the vectors, as it is here, then we know that the relative frequencies of the kmers are the same and therefore should not be penalized. However, for Euclidean distances, this is penalized.

Now, we can note that this is not penalized in both cosine similarity and Pearson correlation. For cosine similarity this is because the two vectors are in the exact same direction. For Pearson correlation this is because the correlation between the variables matches. In fact, one can show that my metrics I've used have the same derivative and grows at the same rate as each other. i.e. for any change in Pearson correlation distance metric, we would expect similar changes in the cosine similarity distance metric. In fact, looking at the previous figures, there is very little discrepancy between the two models. Therefore this boils down to how interpretable the two metrics are.

I argue that the cosine similarity is a more interpretable metric. Cosine similarity simply measures the angle between two vectors. This therefore measures the relative changes in components of both vectors. For frequency, we can imagine that this means how much the relative frequency matches with each other. This relative frequency can then be used as a heuristic for how similar two species are.

However, for the Pearson correlation, we are assuming that the various frequencies act essentially as independent samples of a distribution that might be correlated with another distribution. It is not exactly clear what $\rho = 0$ and its difference from $\rho = -1$, and how much this negative correlation should be taken into account as dissimilarity.

Therefore for its interpretability and usage as a metric, I believe that the cosine similarity should be the best metric to use.